

Using the Standard Template Library (STL)

The Standard Template library

- developed by Alexander Stepanov
- part of the official C++ standard in 97
- most notable example of generic programming and abstractness
- very efficient
- This lecture is about using and applying the STL in your own codes.

The Standard Template library

- **containers**: wrappers for data, can be sequential (`vector`, `list`, `deque`), associative (`map`, `set`, `hash`) or adaptive (`stack`, `queue` — C++11)
- **iterators**: are the major feature that allow the generality of the STL. There are 5 types: *input*, *output*, *forward*, *bidirectional* and *random access* iterators
- **algorithms**: typically searching and sorting (`binary_search`, `lower_bound`, ...) and often requires a custom `operator <` which must guarantee a strict weak ordering
- **functions**: certain classes can overload the function call `operator()`. Such instances are known as *functors*. A typical example is a *predicate* (= a boolean valued function) used eg in `find_if`: it takes a unary predicate that operates on the elements of a sequence
- **allocators**: used for dynamic memory allocation when the size of the containers changes. This will all occur internally and be of no concern in this lecture, but you can use the STL allocators to allocate memory for your own objects

$O(N)$ notation

- when analyzing data structures and algorithms we are interested in
 - how much CPU time is required?
 - how much memory is required?
 - how big is the bandwidth (transfer of data)?
- the crucial question is how these numbers **scale** when increasing (eg doubling) the input size: constant, polynomially or exponentially? A difference between worst-case, best-case and typical scenario can also be made

$O(N)$ notation

- example : addition as taught in primary school

$$\begin{array}{r} 101 \\ 215 \\ + 846 \\ \hline 1061 \end{array}$$

size of input : n digits per input number
however there are also n carry digits
so we have to perform $2n$ additions

time required: $T(n) = 2n$

summing is a linear function in the length of the input

$O(N)$ notation

- example : multiplication as taught in primary school

$$\begin{array}{r} 151 \\ \times 175 \\ \hline 755 \\ 1057 \\ + 151 \\ \hline 26425 \end{array}$$

size of input : n digits per input number


number of multiplications : $n \times n = n^2$

number of additions : also proportional to n^2

time required: $T(n) = n^2 + an^2 = cn^2$

multiplication is a quadratic function in the input


$O(N)$ notation



$O(N)$ notation

definition:

the function $T(n) = \mathcal{O}(g(n))$ when there exist constants c and n_0 such that $T(n) \leq cg(n) \quad \forall n > n_0$



note: mathematicians also use the little-o notation: $T(n) = o(g(n))$ which holds when there exists a constant n_0 such that $T(n) \leq \epsilon g(n) \quad \forall n > n_0$ for *arbitrary* ϵ . It implies that the ratio $T(n) / g(n)$ vanishes in the large n limit

$O(N)$ notation

examples:

$$T(n) = n^2 = \mathcal{O}(n^2)$$

$$T(n) = n^2 + 2n = \mathcal{O}(n^2)$$

$$T(n) = n^2 + \log(n) = \mathcal{O}(n^2)$$

$$T(n) = n^2 + 2n = \mathcal{O}(n^3)$$

since big-O only introduces an upper bound, computer scientists introduce the following 2 definitions:

complexity analysis

definition (for best case analysis):

the function $T(n) = \Omega(g(n))$ when there exist constants c and n_0 such that $T(n) \geq cg(n) \quad \forall n > n_0$

definition (if best and worst case are the same):

the function $T(n) = \theta(g(n))$ when
 $T(n) = \mathcal{O}(g(n))$ and $T(n) = \Omega(g(n))$

quite often, the big-O notation is used where really the θ -notation is meant!

assuming 10 GFlop (~ 2.5 GHz processor) and an operation takes about 0.1 ns

complexity	N=10	10²	10³	10⁴	10⁵	10⁶
1	0.1 ns	0.1 ns	0.1 ns	0.1 ns	0.1 ns	0.1 ns
log N	0.3 ns	0.7 ns	1.0 ns	1.3 ns	1.7 ns	2.0 ns
N	1 ns	10 ns	100 ns	1 μs	10 μs	0.1 ms
N logN	3.3 ns	66.4 ns	1 μs	13 μs	0.17 ms	2 ms
N²	10 ns	1 μs	0.1 ms	10 ms	1 s	1.7 min
N³	0.1 μs	0.1 ms	0.1 s	1.7 min	> 1 day	> 3 years
2^N	0.1 μs	10 ¹⁴ years	10 ²⁸⁵ years			

A little quiz

consider two matrices of size $N \times N$. What is the complexity of their multiplication?

- A. quadratic $\theta(N^2)$
- B. cubic $\theta(N^3)$
- C. sub-cubic $\theta(N^p), p < 3$
- D. exponential $\theta(c^N)$

and what is the complexity of a matrix-vector multiplication?

is there a difference between the computer-time and storage complexity in any of these cases?

A little quiz

we compute the fibonacci numbers recursively

```
double fib(unsigned int n) {  
    if (n == 1) return 1;  
    if (n == 0) return 0;  
    return (fib(n-1) + fib(n-2));  
}
```

what is the complexity of this algorithm?

- | | |
|----------------|--------------------|
| A. constant | $\theta(1)$ |
| B. linear | $\theta(N)$ |
| C. quadratic | $\theta(N^2)$ |
| D. exponential | $\mathcal{O}(c^N)$ |

is there a faster way?

A little quiz

consider the (unnormalized) discrete Fourier transform of data $f(j)$ of length N

$$F(k) = \sum_{j=0}^{N-1} e^{2\pi \sqrt{-1} j k / N} f(j)$$

what is the complexity of computing Fourier transforms?




- | | |
|--------------|---------------|
| A. constant | $\theta(1)$ |
| B. linear | $\theta(N)$ |
| C. quadratic | $\theta(N^2)$ |
| D. other | |

Something to think about

- what is the complexity of finding a number in an array of numbers?
- what is the complexity of finding a name in a telephone book?
- what is the complexity of sorting all student names alphabetically?

Something to think about


Eulerian circuit problem



- 7 bridges of Königsberg
- is there a roundtrip that crosses each bridge (edge) exactly one time?
- more general: what is the complexity class of the best algorithm you can think of for N bridges?

Something to think about

Hamilton cycle problem



- is there a path that crosses each vertex exactly once?
- more general: what is the complexity of the best algorithm you can think of for N vertices?

Standard containers

- pair : a **tuple** of 2 elements (general tuples exist since C++11)
- array structures : **vector**, deque, valarray
- list : **list**
- tree : **map**, multimap, set, multiset
- queue : queue, priority_queue, stack
- **string** and wstring

Since C++11 there is support for a **hash table** implemented in the unordered_map and the unordered_set (see also the lecture on C++11)

pair

- a container for a pair of elements of possibly distinct type

```
template <class T1, class T2> class pair {
public:
    T1 first;
    T2 second;
    pair(const T1& f, const T2& s) : first(f), second(s) {}
};
```

an important function is `make_pair`:

```
template <class T1, class T2>
pair<T1, T2> make_pair (T1 x, T2 y)
{
    return ( pair<T1, T2>(x, y) );
}
```

example modified from cplusplus.com

```
// make_pair example
#include <utility>      // std::pair
#include <iostream>     // std::cout

int main () {
    std::pair <int, int> foo;
    std::pair <int, int> bar;
    std::pair <int, char> baz;

    foo = std::make_pair (10, 20);
    bar = std::make_pair (10.5, 'A'); // ok: implicit conversion from pair<double, char>
    baz = std::make_pair (10.5, 'A');

    std::cout << "foo: " << foo.first << ", " << foo.second << '\n';
    std::cout << "bar: " << bar.first << ", " << bar.second << '\n';
    std::cout << "baz: " << baz.first << ", " << baz.second << '\n';


    return 0;
}
```

String and wstring classes

is usually not so important for numerical work
has partly been addressed in the first lecture;
will however be addressed again when
discussing streams
read your C++ book

stack data structure: LIFO


- **last-in first-out** or last-come first-serve



example : piled-up trays at the mensa



- linear data structure
- single pointer to top of the stack needed
- **peek** : retrieve top element without removing it



- pushing an element to the top is an $O(1)$ operation
- popping the top (=last) element is an $O(1)$ operation
- **peeking** the top element is an $O(1)$ operation

stack

- implementation of a LIFO structure
- functions : empty, size, top, push, and pop

```
// example from cplusplus.com : stack::push/pop
#include <iostream>          // std::cout
#include <stack>              // std::stack

int main ()
{
    std::stack<int> mystack;


    for (int i=0; i<5; ++i) mystack.push(i);

    std::cout << "Popping out elements...";
    while (!mystack.empty())
    {
        std::cout << ' ' << mystack.top();
        mystack.pop();
    }
    std::cout << '\n';

    return 0;
}
```

queue data structure: FIFO

- **first-in first-out** or first-come first-serve




example : waiting line



- linear data structure
- single pointer to top of the stack needed
- **peek** : retrieve top element without removing it

- pushing an element to the top is an $O(1)$ operation
- popping the first element is an $O(1)$ operation
- **peeking** the first and last element is an $O(1)$ operation



queue

- implementation of a FIFO structure
- functions : empty, size, front, back, push, and pop

```
#include <iostream>           // std::cout
#include <queue>               // std::queue

int main ()
{
    std::queue<int> myqueue;

    myqueue.push(100);
    myqueue.push(20);
    myqueue.push(10);

    myqueue.front() -= myqueue.back();    // 100 - 10 = 90

    std::cout << "myqueue.front() is now " << myqueue.front() << '\n';
    for (int i=myqueue.size()-1; i >=0; i--) {
        std::cout << myqueue.front() << "\n";
        myqueue.pop();
    }
    return 0;
}
```


- there also exists a priority_queue : a queue in which the first element is always the greatest element following a strict order (see documentation)

vector


<http://www.cplusplus.com/reference/vector/vector/>

- is a sequence container that can change in size
- use **contiguous** storage locations for their elements
- size can change dynamically; vectors may allocate extra storage, *so the capacity may be larger than its size*
- fast lookup of its elements
- relatively efficient adding/removing elements from the end
- slow in inserting/removing elements at other positions

slow $O(N)$ insertion and removal:
when inserting an element in the
middle, all the elements to the
right must be copied and moved



amortized $O(1)$ for insertion/
removal at the end (spare
capacity!)



see [ex_vector.cpp](#) for its usage and more concepts

the `std::vector` implements what computer scientists understand under the array (aka vector) data structure
note that the specialization `vector<bool>` may behave unexpectedly, see <http://www.cplusplus.com/reference/vector/vector-bool/>


deque

- pronounced “deck”
- is a **D**ouble-**E**nded **Q**ueue
- efficient insertion and deletion at the beginning and at the end of the data sequence
- does NOT store data in contiguous memory (hence pointer arithmetic is undefined)
- has random access iterators

list

<http://www.cplusplus.com/reference/list/list/>

- is a sequence container that can change in size;
[not contiguous](#) in memory
- O(1) insertion and removal everywhere in the list
- doubly-linked (the forward list is singly-linked)
- lacks direct access to the elements by its position




see [ex_list.cpp](#) for its usage and more concepts

Tree structures

- an array needs
 - $O(1)$ for random access
 - $O(N)$ for arbitrary insertions and removals
 - $O(N)$ for searches
 - $O(\log N)$ for searches in a sorted array
- A list needs
 - $O(1)$ for arbitrary insertion and removal
 - $O(N)$ for random access and searches
- what if both operations need to be fast? Use tree data structure
 - $O(\log N)$ for arbitrary insertion and removal
 - $O(\log N)$ for random access and searches


Tree structures

- A tree is an undirected graph in which any 2 vertices are connected by exactly one path
- in computer science rooted trees are important: one designated vertex from which all edges point away




example of a tree
there are n vertices, $n-1$ edges

- A binary tree : every node has (at most) 2 child nodes, the right child node is larger than the parent and the left child node



binary tree


- a tree of height n can store $N = 2^n - 1$ elements
- if the tree is perfectly balanced, then a search can take no more than n operations



- trees can however become unbalanced:
imagine inserting elements from a sorted array:


this is problematic because all operations become $O(N)$ instead of $O(\log N)$

solution : rebalancing methods (eg rotations), or use self-balancing trees



binary search tree

- in a binary search tree all elements on a right subtree must be larger than those on any parent node along the subtree



is not a BST

- search : start at root, compare value, go left or right etc
- insert: start with search, then insert new key-value pair as new leaf
- delete: a bit more complicated (the node to be deleted may have dependencies) — see textbooks

see textbooks on computer science for more tree structures (red-black tree, 2-3-4 tree, ...)

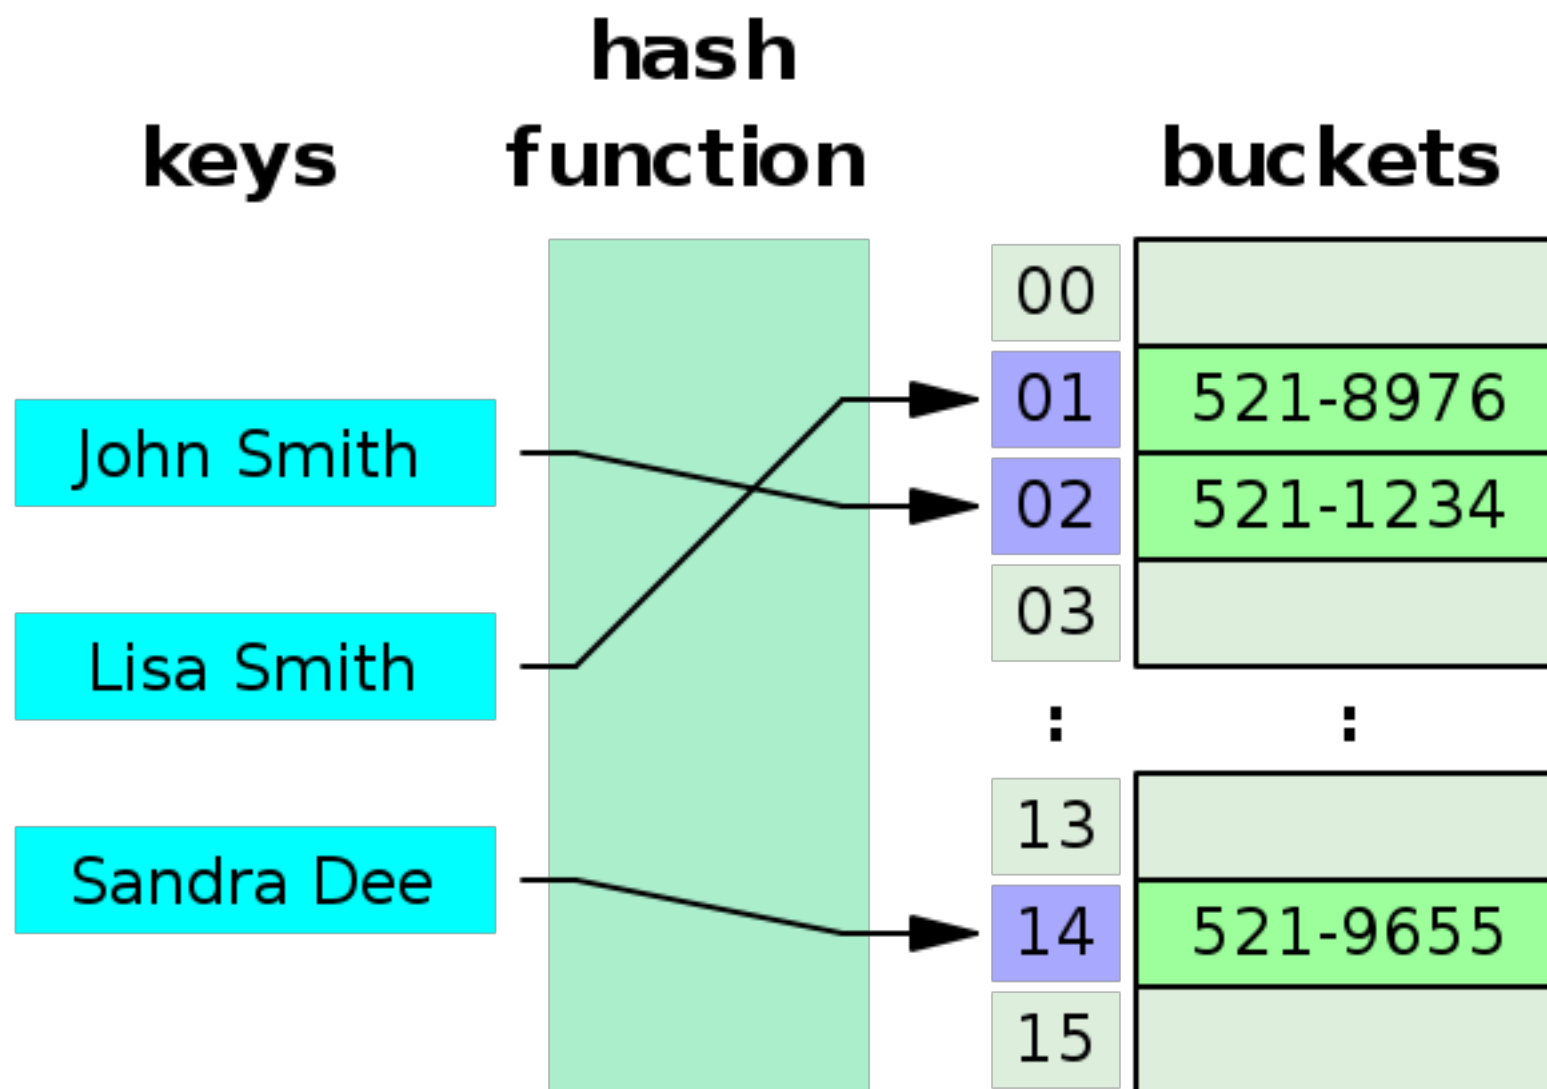
map

- are associative containers storing elements as pairs `<key, value>`. The keys are usually used to sort (following a strict weak ordering criterion) and identify the elements *uniquely*
- are implemented as binary search trees
- has the usual iterators: `begin`, `end`, `rbegin`, `rend`, ...
- has functions : `insert`, `erase`, `clear`, `empty`, `size`, ...
- has operator[] to access the elements — `map[key]` returns `map_value`
- has operations: `find`, `count`, `lower_bound`, `upper_bound`, `equal_range`

map example

- see the example [ex_map.cpp](#) in this week's programs for usage
- [multimap](#) : key does not have to be unique (see the documentation)
- [set](#) : ordering is based on values only; value must be unique (see the documentation)
- [multiset](#) : like set, but value must not be unique
- do not be confused: the [unordered_map](#), [unordered_multimap](#), [unordered_set](#), [unordered_multiset](#) (C++11) : is not ordered internally, but uses a *hash* function

hash table



```
index = f(key, array_size)
```

often done in two steps:

```
hash = hashfunc(key)
index = hash % array_size
```

a good hash gives a good idea where the index can be found

Hash table


Type	Unordered associative array
Invented	1953

Time complexity		
	Average	Worst case
Space	$O(n)^{[1]}$	$O(n)$
Search	$O(1)$	$O(n)$
Insert	$O(1)$	$O(n)$
Delete	$O(1)$	$O(n)$

hash table

- typical situation in physics: diagonalization. The Hamiltonian is stored as a sparse matrix

most elements are zero!



how to store? many options, for instance a coordinate list COO with tuples (row, column, value) sorted first by the row index, then by the column index

when H acts on a vector, we need to know the *index* of the final state

how does a hashing function work?

- consider the exact diagonalization of a spin-1/2 system. Let there be L sites and consider the magnetization sector $M = 0$ — ie, we assume that the S^z operator commutes with the Hamiltonian.
- without conservation of the magnetization there would be 2^L states
- in the submanifold there are however fewer states: $\frac{L!}{(L/2)!(L/2)!}$
- states can be represented by bits:
 $\uparrow \rightarrow 1 \qquad \downarrow \rightarrow 0$
- a state can hence be represented as an integer (this is a bijection):
 $[1, 0, 0, 1] = 9$
- this allows for fast bit operations via AND, OR and XOR

how does a hashing function work?

- for $L = 4$, there are 16 states in total and 6 states with magnetization 0

$$[0, 0, 1, 1] = 3$$

$$[0, 1, 0, 1] = 5$$

$$[0, 1, 1, 0] = 6$$

$$[1, 0, 0, 1] = 9$$

$$[1, 0, 1, 0] = 10$$

$$[1, 1, 0, 0] = 12$$

we can iterate over all 2^L states and check the magnetization; alternatively, we start with the state with lowest possible integer representation and attempt to move iteratively the right-most 1-bits to the left until we reach another 1 or the end of the sequence (see exercises)

- we only want to store these 6 valid states. The problem now is to find the index of a valid state. Eg, given $[1001] = 9$ we need to know that it has index 3.
- solution nr 1: use a binary search tree (`std::set` or `std::map`) or a sorted vector. In this case the lookup is logarithmic in the the number of states ($O(\log N)$) which in most cases is fast enough.

how does a hashing function work?

- solution nr 2: with a hashtable we can provide lookup with $O(1)$. This can be done with `std::unordered_multimap`. To show how hashing works, we will implement our own hash function

state	index	hash
$[0, 0, 1, 1] = 3$	0	3
$[0, 1, 0, 1] = 5$	1	5
$[0, 1, 1, 0] = 6$	2	0
$[1, 0, 0, 1] = 9$	3	3
$[1, 0, 1, 0] = 10$	4	4
$[1, 1, 0, 0] = 12$	5	0

- a simple, though not the optimal, hash function is $\text{hash} = (\text{state} \% 6)$

how does a hashing function work?

- we can now store two arrays. The first one consists of $\langle \text{key}, \text{value} \rangle$ pairs where key is the result of the hash function and value the decimal representation of the state:

$[\langle 0, 2 \rangle, \langle 0, 5 \rangle, \langle 3, 0 \rangle, \langle 3, 3 \rangle, \langle 4, 4 \rangle, \langle 5, 1 \rangle]$

- one sees that the keys are not unique because our hash function is not perfect. In the above example there are 2 collisions
- the above array does not need to be ordered, but entries with the same keys ('buckets') must be consecutive.
- the second array is the index array H:

$[0, \text{nan}, \text{nan}, 2, 4, 5]$

- $H[n] = p$ it tells us the location of the first entry in the $\langle \text{key}, \text{value} \rangle$ array where the key is p.
- nan means here that 1 and 2 cannot be valid results of our hash function over the chosen Hilbert space

example

- we wish to know the index of $[1, 0, 1, 0]$

decimal : 10 ; hash result : 4.

result of index array : 4. The corresponding value to this key is unique, namely 4; ie the index of $[1,0,1,0]$ is 4.

- we wish to know the index of $[1, 0, 0, 1]$

decimal : 9 ; hash result : 3.

result of index array : 2. The corresponding value to this key is not unique, both 0 and 3 are valid. Then we need to check explicitly which state we want.

After comparing the 2 values with the initial state, we deduce that its index is 3

- for large Hilbert spaces and very good hash functions this can produce a lookup with $O(1)$ complexity.
- it functions best when the manifold of states remains invariant (ie inserts and removes are seldom)
- a good hash function is usually given by $(\cdot \% p)$ with p the smallest prime number larger than the size of the Hilbert space.
- in the exercises you will work out this example

unordered_map

```
// unordered_map -- requires C++11
#include <iostream>
#include <string>
#include <unordered_map>

int main ()
{
    std::unordered_map<std::string,int> bachelor_students, all_students;

    std::pair<std::string,int> exchange_student ("Fabricio", 2085);
    all_students.insert(exchange_student);
    bachelor_students.insert(std::make_pair<std::string,int>("Ann", 2183));
    bachelor_students.insert(std::make_pair<std::string,int>("Tom", 2184));
    all_students.insert(bachelor_students.begin(), bachelor_students.end());

    std::cout << "all_students contains " << std::endl;
    for (std::unordered_map<std::string,int>::iterator it = all_students.begin(); it != all_students.end(); ++it) {
        std::cout << it->first << " : " << it->second << std::endl;
    }

    std::cout << std::endl;

    std::unordered_map<std::string,int>::hasher hfun = all_students.hash_function();
    for (std::unordered_map<std::string,int>::iterator it = all_students.begin(); it != all_students.end(); ++it) {
        std::cout << hfun(it->first) << std::endl;
    }

    return 0;
}
```

STL algorithms

- The STL defines a huge number of algorithms that work on almost all containers in `<algorithm>` :

see <http://en.cppreference.com/w/cpp/algorithm> for a complete list

- a few examples:
 - `count` : counts how often a predicate is true over a sequence

```
int num_items1 = std::count(v.begin(), v.end(), target1);
```

- `find` : searches for an element equal to value

```
std::vector<int>::iterator it = std::find(v.begin(), v.end(), 3);
```

- `copy` : copies a range

```
std::vector<int> to_vector;  
std::copy(from_vector.begin(), from_vector.end(),  
          std::back_inserter(to_vector));  
  
std::cout << "to_vector contains: ";  
  
std::copy(to_vector.begin(), to_vector.end(),  
          std::ostream_iterator<int>(std::cout, " "));  
std::cout << '\n';
```

STL algorithms

- `swap` : swaps the contents of 2 elements

```
int a = 5, b = 3;  
std::swap(a,b);
```

- `unique` : removes duplicate elements from a sequence

```
int myints[] = {0,7,5,6,1,2,2,3,4,7,4,3,3,3};  
std::vector<int> v (myints, myints + sizeof(myints) / sizeof(int) );  
std::sort(v.begin(), v.end()); // 0 1 2 2 3 3 3 3 4 4 5 6 7 7  
std::vector<int>::iterator last = std::unique(v.begin(), v.end()); // 0 1 2 3 4 5 6 7 x x x x x x  
v.erase(last, v.end()); // remove indeterminate elements
```

- `sort` : sorts a range in ascending order
- `lower_bound` : (on sorted range): returns an iterator to the first element not less than the given value
- `upper_bound` : (on sorted range): returns an iterator to the first element greater than the given value

STL algorithms

- `max` : returns the greater of two given values
- `max_element` : returns the largest element in a range
- `min` : returns the smaller of two given values
- `min_element` : returns the smallest element in a range
- `accumulate` : sums up a range of elements
- `inner_product` : returns the inner product
- `next_permutation` :

```
int main()
{
    std::string s = "aba";
    std::sort(s.begin(), s.end());
    do {
        std::cout << s << '\n';
    } while(std::next_permutation(s.begin(), s.end()));
}
```

→
aab
aba
baa

homework

- write a simple function which returns the minimum of 2 integers
- write a simple function which returns the minimum of 2 doubles
- write a simple function which returns the minimum of 2 floats